

16. Data Type Control (NumPy)

1. Topic Overview

What this topic is

Data Type Control means controlling **how numbers are stored in memory** inside NumPy arrays.

Each NumPy array has **one fixed data type (dtype)**.

Why it exists

Computers store data in memory.

Different data types use different memory sizes.

Correct dtype gives:

- Less memory usage
- Faster computation
- Correct numerical results

One real-world analogy

Think of containers:

- Small bottle for water
 - Big tank for fuel
- Using the wrong container wastes space or overflows.

2. Core Theory (Deep but Clear)

NumPy arrays are:

- Homogeneous (one dtype only)
- Stored in contiguous memory

- Optimized for math operations

Internal NumPy view

Each array has:

- **shape** → how many rows and columns
- **dtype** → how each element is stored
- **itemsize** → bytes per element
- **memory block** → continuous memory

If dtype is wrong:

- Memory waste happens
- Precision loss happens
- ML models behave badly

Sub-topics in Data Type Control

1. `dtype`
2. Common numeric dtypes
3. `astype()`
4. `itemsize`
5. Type inference
6. Overflow and precision issues

3. Syntax & Examples

3.1 `dtype`

What it is

Defines the type of each element in the array.

Basic syntax

```
import numpy as np  
  
arr = np.array([1, 2, 3], dtype=np.int32)
```

Example 1

```
arr = np.array([1, 2, 3])  
print(arr.dtype)
```

Output

```
int64
```

Explanation

- NumPy inferred type automatically
- Uses platform default integer

Example 2

```
arr = np.array([1, 2, 3], dtype=np.float32)  
print(arr)  
print(arr.dtype)
```

Output

```
[1. 2. 3.]  
float32
```

Explanation

- Integers converted to floats
- Each value stored in 4 bytes

3.2 Common Numeric dtypes

Type	Meaning	Bytes
int8	small integers	1
int32	standard integers	4
int64	large integers	8
float32	decimal numbers	4
float64	high precision	8

Example

```
arr = np.array([1, 2, 3], dtype=np.int8)
print(arr)
print(arr.itemsize)
```

Output

```
[1 2 3]
1
```

Explanation

- Each element uses 1 byte
- Very memory efficient
- Risk of overflow

3.3 astype()

What it does

Changes dtype of an existing array.

Creates a **new copy**.

Syntax

```
new_arr = arr.astype(np.float32)
```

Example 1

```
arr = np.array([1, 2, 3])
new_arr = arr.astype(np.float32)

print(arr.dtype)
print(new_arr.dtype)
```

Output

```
int64
float32
```

Explanation

- Original array unchanged
- New array created

Example 2

```
arr = np.array([1.7, 2.3, 3.9])
new_arr = arr.astype(int)

print(new_arr)
```

Output

```
[1 2 3]
```

Explanation

- Decimal part removed
- No rounding
- Dangerous in ML pipelines

3.4 itemsize

What it is

Memory used by **one element**.

Example

```
arr = np.array([1, 2, 3], dtype=np.float64)
print(arr.itemsize)
```

Output

8

Explanation

- Each value uses 8 bytes
- Large arrays become heavy

3.5 Type Inference

NumPy picks dtype based on input.

Example

```
arr = np.array([1, 2.5, 3])
print(arr.dtype)
```

Output

float64

Explanation

- Mixed types
- NumPy chooses safest type

3.6 Overflow and Precision

Example (Overflow)

```
arr = np.array([120, 121], dtype=np.int8)
print(arr + 10)
```

Output

```
[ -126 -125]
```

Explanation

- int8 range exceeded
- Values wrapped
- Silent error

4. Why This Matters in Data Science

Data cleaning

- Wrong dtype breaks missing value handling
- Strings stored as numbers cause errors

Feature engineering

- Categorical encoding needs correct dtype
- Scaling fails if integers are used

Model input preparation

- Deep learning models expect float32
- Wrong dtype slows GPU training

ML / DL pipelines

- Overflow gives wrong gradients
- Precision loss ruins model accuracy

If you don't understand this

- Models train but give wrong results
- Memory usage explodes
- Bugs become invisible

5. Common Mistakes (VERY IMPORTANT)

1. Using default dtype blindly

Cause: Trusting NumPy too much

Fix: Always check `dtype`

2. Converting floats to int unknowingly

Cause: `astype(int)`

Fix: Convert only after validation

3. Ignoring overflow

Cause: Small integer types

Fix: Use `int32` or `int64`

4. Mixing numeric and string data

Cause: Dirty datasets

Fix: Clean before array creation

5. Using float64 everywhere

Cause: Habit

Fix: Use `float32` for ML

6. Performance & Best Practices

When this is fast

- Correct dtype
- Contiguous memory
- Vectorized operations

When this is slow

- Unnecessary type conversions
- Large float64 arrays
- Repeated `astype()` calls

Memory warnings

- float64 doubles memory vs float32
- int8 can overflow silently

7. Practice Problems

Easy (5)

1. Create an array of 100 zeros using `float32`
2. Check dtype and itemsize of an integer array
3. Convert a float array to int
4. Identify dtype chosen by NumPy for mixed input
5. Create an array with `int16` and print memory usage

Medium (7)

1. Convert dataset values to `float32` for ML input
2. Detect overflow in small integer arrays
3. Reduce memory of a large array safely
4. Compare memory of `float32` vs `float64`

5. Fix dtype issues in a noisy dataset
6. Prepare model input array with correct dtype
7. Convert categorical labels to integers

Hard (5)

1. Debug wrong predictions caused by dtype
2. Fix precision loss in probability outputs
3. Optimize memory for million-row dataset
4. Detect silent overflow in pipeline
5. Prepare NumPy arrays for GPU training

Industry-Level Tasks (3)

1. Convert raw CSV data into ML-ready NumPy arrays
Input: mixed types
Output: optimized dtype arrays
2. Memory optimization for real-time inference system
Constraint: limited RAM
3. Fix training instability caused by dtype mismatch
Scenario: deep learning model

8. Mini Checklist

- NumPy arrays have one dtype only
- dtype controls memory and speed
- Always check `dtype` and `itemsize`
- `astype()` creates a copy
- `float32` is standard for ML
- Small integers can overflow
- Wrong dtype breaks models silently

16 (Part 2). Advanced Data Type Control (Mandatory for Data Science)

1. Topic Overview

What this topic is

These are **advanced NumPy dtypes** used for:

- dates
- time differences
- boolean masks
- mixed structured data
- raw Python objects

Why it exists

Real datasets are not only numbers.

They contain:

- dates
- categories
- flags
- mixed columns

NumPy needs special dtypes to handle them safely.

One real-world analogy

Excel columns.

Date column is not same as number column.

Each needs different rules.

2. Core Theory (Deep but Clear)

NumPy dtype system is larger than `int` and `float`.

Important advanced dtypes:

1. `bool`
2. `datetime64`
3. `timedelta64`
4. `object`
5. Structured dtypes
6. `view()` vs `astype()`

If you ignore these:

- Time-series analysis breaks
- Masking becomes slow
- Pipelines crash silently

3. Syntax & Examples

3.1 Boolean dtype (`bool`)

What it is

Stores `True` or `False`.

Uses **1 byte per value**.

Used heavily for filtering.

Syntax

```
np.array(data, dtype=bool)
```

Example 1

```
import numpy as np

arr = np.array([1, 0, 5, 0], dtype=bool)
print(arr)
print(arr.dtype)
```

Output

```
[ True False  True False]
bool
```

Explanation

- Non-zero becomes True
- Zero becomes False

Example 2 (Masking)

```
data = np.array([10, 20, 30, 40])
mask = data > 25

print(mask)
print(data[mask])
```

Output

```
[False False  True  True]
[30 40]
```

Explanation

- Boolean array used as filter
- Core Data Science operation

3.2 datetime64

What it is

Stores dates and timestamps.

Much faster than Python `datetime`.

Syntax

```
np.array(dates, dtype='datetime64')
```

Example 1

```
dates = np.array(['2024-01-01', '2024-01-10'], dtype='datetime64')
print(dates)
print(dates.dtype)
```

Output

```
['2024-01-01' '2024-01-10']
datetime64[D]
```

Explanation

- Stored as integers internally
- Unit is days

Example 2

```
print(dates[1] - dates[0])
```

Output

```
9 days
```

Explanation

- Vectorized date math
- No loops needed

3.3 timedelta64

What it is

Stores time difference.

Used with `datetime64`.

Example

```
delta = np.array([1, 3, 7], dtype='timedelta64[D]')
print(delta)
```

Output

```
[1 3 7]
```

Explanation

- Represents day differences
- Used in time-based features

3.4 object dtype

What it is

Stores **raw Python objects**.

Very slow.

Avoid unless required.

Example

```
arr = np.array(['apple', 10, 3.5])
print(arr.dtype)
```

Output

object

Explanation

- Mixed types force object dtype
- No vectorization
- Memory heavy

3.5 Structured dtypes

What it is

Multiple fields per element.

Like rows with fixed schema.

Example

```
dtype = [('id', 'int32'), ('score', 'float32')]
data = np.array([(1, 90.5), (2, 85.0)], dtype=dtype)

print(data)
print(data['score'])
```

Output

```
[(1, 90.5) (2, 85. )]
[90.5 85. ]
```

Explanation

- Each element has structure

- Used in low-level pipelines

3.6 view() vs astype() (CRITICAL)

astype()

- Changes dtype
- Creates copy

view()

- Reinterprets memory
- No copy

Example

```
arr = np.array([1, 2, 3], dtype=np.int32)

view_arr = arr.view(np.float32)
print(view_arr)
```

Output

```
[1.4012985e-45 2.8025969e-45 4.2038954e-45]
```

Explanation

- Same memory
- Interpreted differently
- Dangerous if misunderstood

4. Why This Matters in Data Science

Data cleaning

- Boolean masks remove bad rows
- Object dtype signals dirty data

Feature engineering

- Date differences create time features
- Flags stored as bool

Model input preparation

- Object dtype breaks ML models
- Dates must be numeric or encoded

ML / DL pipelines

- Time-based splits rely on datetime
- Wrong dtype causes silent failure

If you skip this

- Time series models fail
- Filtering becomes slow
- Pipelines become unstable

5. Common Mistakes (VERY IMPORTANT)

1. Leaving strings as `object`

Cause: Direct CSV load

Fix: Convert early

2. Using Python `datetime` instead of `datetime64`

Cause: Habit

Fix: Use NumPy `datetime`

3. Using `view()` without understanding memory

Cause: Curiosity

Fix: Prefer `astype()`

4. Boolean masks with wrong shape

Cause: Shape mismatch

Fix: Always check shape

5. Structured dtypes for ML inputs

Cause: Overengineering

Fix: Flatten before modeling

6. Performance & Best Practices

Fast

- bool masks
- datetime64 math
- numeric dtypes

Slow

- object dtype
- Python loops
- mixed-type arrays

Warnings

- object dtype disables vectorization
- `view()` can corrupt logic
- datetime unit mismatch causes bugs

7. Practice Problems

Easy (5)

1. Create a boolean mask for values > mean

2. Convert string dates to `datetime64`
3. Subtract two date arrays
4. Detect object dtype in an array
5. Create a boolean array from integers

Medium (7)

1. Filter dataset using boolean masks
2. Create time-gap features from dates
3. Convert object arrays to numeric safely
4. Identify slow operations caused by dtype
5. Fix shape mismatch in boolean indexing
6. Encode date differences as integers
7. Remove object dtype from ML inputs

Hard (5)

1. Debug model crash caused by object dtype
2. Optimize time-series feature generation
3. Replace Python datetime with NumPy datetime
4. Detect memory misuse from `view()`
5. Clean mixed-type NumPy arrays

Industry-Level Tasks (3)

1. Build time-based features for churn prediction
Input: user signup dates
Output: numeric features
2. Optimize filtering on 10M rows using boolean masks
3. Fix production pipeline slowed by object dtype arrays

8. Mini Checklist

- Boolean masks are core to filtering
- datetime64 is mandatory for time data
- object dtype is dangerous
- Structured dtypes are not for ML models
- `astype()` copies, `view()` does not
- Always remove object dtype before modeling